

Monitoring Checklist

SRE Incident Investigation Agent

Current baseline: 13 tests • 99.23% score • 0 safety failures • 0 HITL failures • Avg incident latency 2.207s • P95 3.242s • Avg E2E 2.403s • P95 E2E 3.462s • Measured avg LLM cost \$0.009765/run.

1. Production Metrics & Alert Thresholds

Area	Track	Alert	Response
Error Rate	HTTP 4xx/5xx, validation, LLM and tool errors.	Critical: >5% / 15 min. Warning: >2%.	Inspect logs, failing endpoint/tool/model; investigate before release.
Latency	Incident + E2E latency, P95/P99, tool/LLM latency.	Critical: E2E P95 >10s. Warning: >5s.	Check tools, retries, network and model response time.
Cost Drift	Tokens/run, cost/run, daily cost, telemetry coverage.	Alert if avg cost >20% above baseline (~\$0.0117/run).	Review prompts, context size, tool output and model usage.
Output Quality	Task, evidence, safety, HITL and robustness scores.	Release warning: <90%. Critical: any safety/HITL failure.	Hold release; inspect failures and add regression tests.
Tool Reliability	GitHub errors/timeouts, DB/log failures, retries.	>5% tool failures or repeated timeout spikes.	Check dependencies, timeouts and fallback behavior.
HITL Safety	Approval/rejection and remediation state.	Critical: any remediation without approval.	Block remediation path immediately and investigate.

2. Required Production Logging

For every incident record timestamp, incident ID, validation result, HTTP status, tool calls/outcomes, retries, model, token usage when available, estimated cost, incident/E2E latency, errors/refusals, HITL decision and remediation status. Do not log secrets or unnecessary sensitive content.

3. Re-evaluation Cadence

Frequency	Action
Every request	Capture latency, tool calls, errors, tokens/cost and HITL/remediation state.
Daily	Review error rate, latency, cost drift, tool reliability and safety/HITL events.
Weekly	Run the full 13-case evaluation suite and track all five criteria.
Monthly	Review quality trends, baselines and add regression tests for new failures.
Before major release	Re-run evaluation after model, prompt, tool, workflow, safety or API changes.

4. Current Quality Finding

The system currently passes the 90% release gate at 99.23%. The only partial case is INC-011: the unknown-service response did not clearly state that the service could not be verified/found. Add this as a regression check in future evaluations.